

```
In [1]: from sklearn.datasets import load_breast_cancer
import pandas as pd
dataset=load_breast_cancer()
df=pd.DataFrame(dataset['data'],columns=dataset['feature_names'])
df['target']=dataset['target']
df.head()
```

Out[1]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	...	worst texture	worst perimeter	worst area	worst smoothness
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	17.33	184.60	2019.0	0.1622
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	23.41	158.80	1956.0	0.1238
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	25.53	152.50	1709.0	0.1444
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	...	26.50	98.87	567.7	0.2098
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	16.67	152.20	1575.0	0.1374

5 rows × 31 columns

```
In [2]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   mean radius                           569 non-null    float64
1   mean texture                           569 non-null    float64
2   mean perimeter                         569 non-null    float64
3   mean area                             569 non-null    float64
4   mean smoothness                       569 non-null    float64
5   mean compactness                      569 non-null    float64
6   mean concavity                        569 non-null    float64
7   mean concave points                   569 non-null    float64
8   mean symmetry                         569 non-null    float64
```

```

9  mean fractal dimension    569 non-null    float64
10 radius error             569 non-null    float64
11 texture error            569 non-null    float64
12 perimeter error          569 non-null    float64
13 area error               569 non-null    float64
14 smoothness error         569 non-null    float64
15 compactness error        569 non-null    float64
16 concavity error          569 non-null    float64
17 concave points error     569 non-null    float64
18 symmetry error           569 non-null    float64
19 fractal dimension error   569 non-null    float64
20 worst radius              569 non-null    float64
21 worst texture             569 non-null    float64
22 worst perimeter           569 non-null    float64
23 worst area                569 non-null    float64
24 worst smoothness         569 non-null    float64
25 worst compactness        569 non-null    float64
26 worst concavity          569 non-null    float64
27 worst concave points     569 non-null    float64
28 worst symmetry           569 non-null    float64
29 worst fractal dimension   569 non-null    float64
30 target                   569 non-null    int32

```

dtypes: float64(30), int32(1)

memory usage: 135.7 KB

```

In [3]: from sklearn.model_selection import train_test_split
data=df.loc[:, df.columns != 'target']
target=df['target']
x_train,x_test,y_train,y_test=train_test_split(data,target,test_size=0.20,random_state=1)

```

```

In [4]: print('train data of x_train is: ',x_train.shape)
print('train data of x_test is: ',x_test.shape)
print('train data of y_train is: ',y_train.shape)
print('train data of y_train is: ',y_test.shape)

```

```

train data of x_train is: (455, 30)
train data of x_test is: (114, 30)
train data of y_train is: (455,)
train data of y_train is: (114,)

```

```

In [5]: print(y_train.value_counts())

```

```
print(y_test.value_counts())
```

```
1    285
0    170
Name: target, dtype: int64
1     72
0     42
Name: target, dtype: int64
```

```
In [6]: # import the class
        from sklearn.linear_model import LogisticRegression
        # instantiate the model (using the default parameters)
        logreg = LogisticRegression()
        # fit the model with data
        logreg.fit(x_train,y_train)
        #
        y_pred=logreg.predict(x_test)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
In [7]: # import the metrics class
        from sklearn import metrics
        cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
        cnf_matrix
```

```
Out[7]: array([[38,  4],
               [ 2, 70]], dtype=int64)
```

```
In [8]: print(metrics.classification_report(y_test,y_pred,digits=3))
```

	precision	recall	f1-score	support
0	0.950	0.905	0.927	42

1	0.946	0.972	0.959	72
accuracy			0.947	114
macro avg	0.948	0.938	0.943	114
weighted avg	0.947	0.947	0.947	114

```
In [9]: import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import *
from sklearn.naive_bayes import GaussianNB

import matplotlib.pyplot as plt
```

```
In [22]: import warnings
warnings.filterwarnings("ignore")
```

```
In [10]: df.head()
```

```
Out[10]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	...	worst texture	worst perimeter	worst area	worst smoothness
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	17.33	184.60	2019.0	0.1622
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	23.41	158.80	1956.0	0.1238
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	25.53	152.50	1709.0	0.1444
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	...	26.50	98.87	567.7	0.2098
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	16.67	152.20	1575.0	0.1374

5 rows × 31 columns

In [11]:

```
clf1 = RandomForestClassifier(random_state=42)
clf2 = SVC(probability=True, random_state=42)
clf3 = LogisticRegression(random_state=42)
clf4 = DecisionTreeClassifier(random_state=42)
clf5 = KNeighborsClassifier()
clf6 = GaussianNB()
clf7 = GradientBoostingClassifier(random_state=42)
```

In [12]:

```
param1 = {}
param1['classifier__n_estimators'] = [10, 50, 100, 250]
param1['classifier__max_depth'] = [5, 10, 20]
param1['classifier__class_weight'] = [None, {0:1,1:5}, {0:1,1:10}, {0:1,1:25}]
param1['classifier'] = [clf1]
param2 = {}
param2['classifier__C'] = [10**-2, 10**-1, 10**0, 10**1, 10**2]
param2['classifier__class_weight'] = [None, {0:1,1:5}, {0:1,1:10}, {0:1,1:25}]
param2['classifier'] = [clf2]
param3 = {}
param3['classifier__C'] = [10**-2, 10**-1, 10**0, 10**1, 10**2]
param3['classifier__penalty'] = ['l1', 'l2']
param3['classifier__class_weight'] = [None, {0:1,1:5}, {0:1,1:10}, {0:1,1:25}]
param3['classifier'] = [clf3]
param4 = {}
param4['classifier__max_depth'] = [5,10,25,None]
param4['classifier__min_samples_split'] = [2,5,10]
param4['classifier__class_weight'] = [None, {0:1,1:5}, {0:1,1:10}, {0:1,1:25}]
param4['classifier'] = [clf4]
param5 = {}
param5['classifier__n_neighbors'] = [2,5,10,25,50]
param5['classifier'] = [clf5]
param6 = {}
# param6['var_smoothing'] = np.logspace(0,-9, num=100)
param6['classifier'] = [clf6]
param7 = {}
param7['classifier__n_estimators'] = [10, 50, 100, 250]
param7['classifier__max_depth'] = [5, 10, 20]
param7['classifier'] = [clf7]
```

In [35]:

```

import matplotlib.pyplot as plt

c=0
for i in (clf1,clf2,clf3,clf4,clf5,clf6,clf7):

    pipeline = Pipeline([('classifier', i)])
    params = [param1, param2, param3, param4, param5, param6, param7]
    gs = GridSearchCV(pipeline, params[c], cv=3, n_jobs=-1, scoring='roc_auc').fit(x_train, y_train)
    # print(gs.best_params_)
    print(i)
    # print(pipeline, params[c])
    c=c+1
    if(i==i):
        y_pred_proba = gs.predict_proba(x_test)[:,:1]
        fpr_k, tpr_k, _k = metrics.roc_curve(y_test, y_pred_proba)
        auc = metrics.roc_auc_score(y_test, y_pred_proba)
        plt.plot(fpr_k,tpr_k,label=str(i).partition("!")[0]+"KNN, auc="+str(auc))
        plt.legend()
        plt.legend(loc=(1.04,0))

    print("Test ROC AUC Score:",roc_auc_score(gs.predict(x_test), y_test))

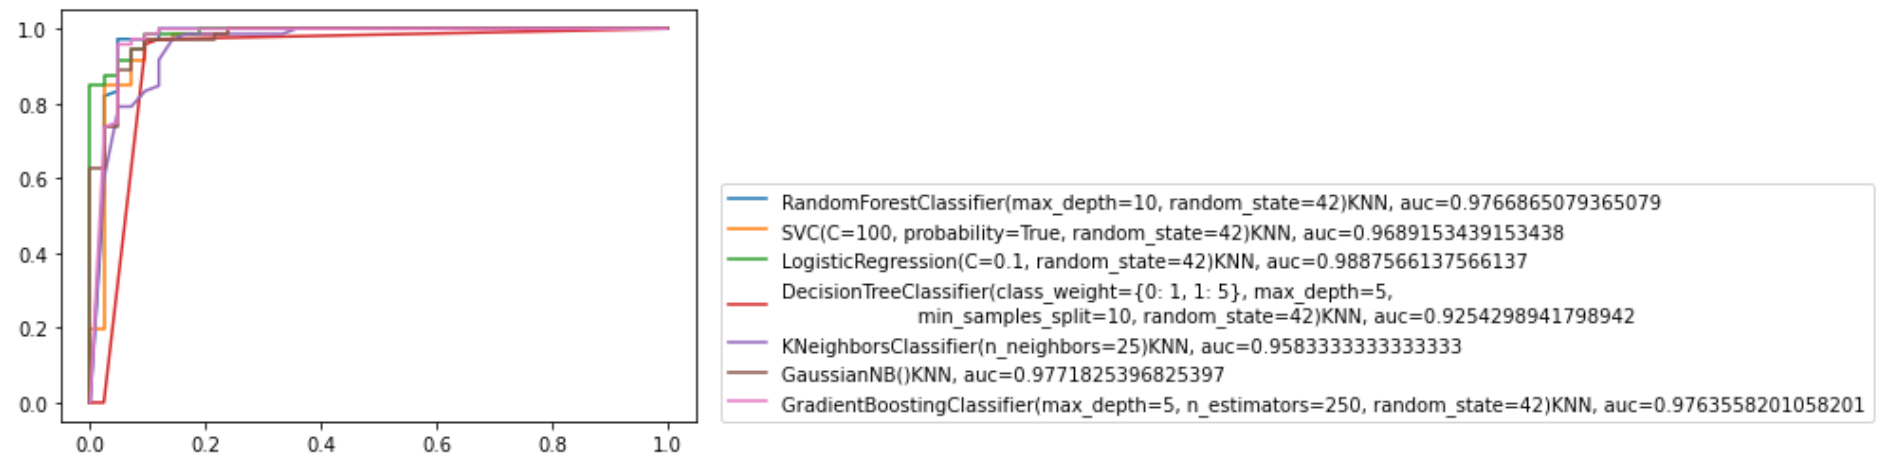
print("\n")

```

```

RandomForestClassifier(max_depth=10, random_state=42)
Test ROC AUC Score: 0.9675324675324675
SVC(C=100, probability=True, random_state=42)
Test ROC AUC Score: 0.9342105263157896
LogisticRegression(C=0.1, random_state=42)
Test ROC AUC Score: 0.9605128205128205
DecisionTreeClassifier(class_weight={0: 1, 1: 5}, max_depth=5,
                        min_samples_split=10, random_state=42)
Test ROC AUC Score: 0.941025641025641
KNeighborsClassifier(n_neighbors=25)
Test ROC AUC Score: 0.9231200897867564
GaussianNB()
Test ROC AUC Score: 0.947972972972973
GradientBoostingClassifier(max_depth=5, n_estimators=250, random_state=42)
Test ROC AUC Score: 0.947972972972973

```



```
In [ ]: print("Best AUC: Logistic Regression 0.98")
```